




# Predicting Home Values with Linear Regression



## Contents of the Jupyter Notebook:

1. Methodology
2. Exploring the Data
3. Splitting, Selecting, and Scaling Data
4. Linear Regression
5. Takeaways



Creating dummy variables:

```
# Create dummy variables
from sklearn import preprocessing
from sklearn.preprocessing import KBinsDiscretizer

waterfront_dummies = pd.get_dummies(df['waterfront'], prefix='wf', drop_first=True)
condition_dummies = pd.get_dummies(df['condition'], prefix='cond', drop_first=True)
```

Binning continuous values:

```
disc_5k = KBinsDiscretizer(encode='onehot-dense', strategy='kmeans')
disc_5 = KBinsDiscretizer(encode='onehot-dense')
disc_3 = KBinsDiscretizer(encode='onehot-dense')

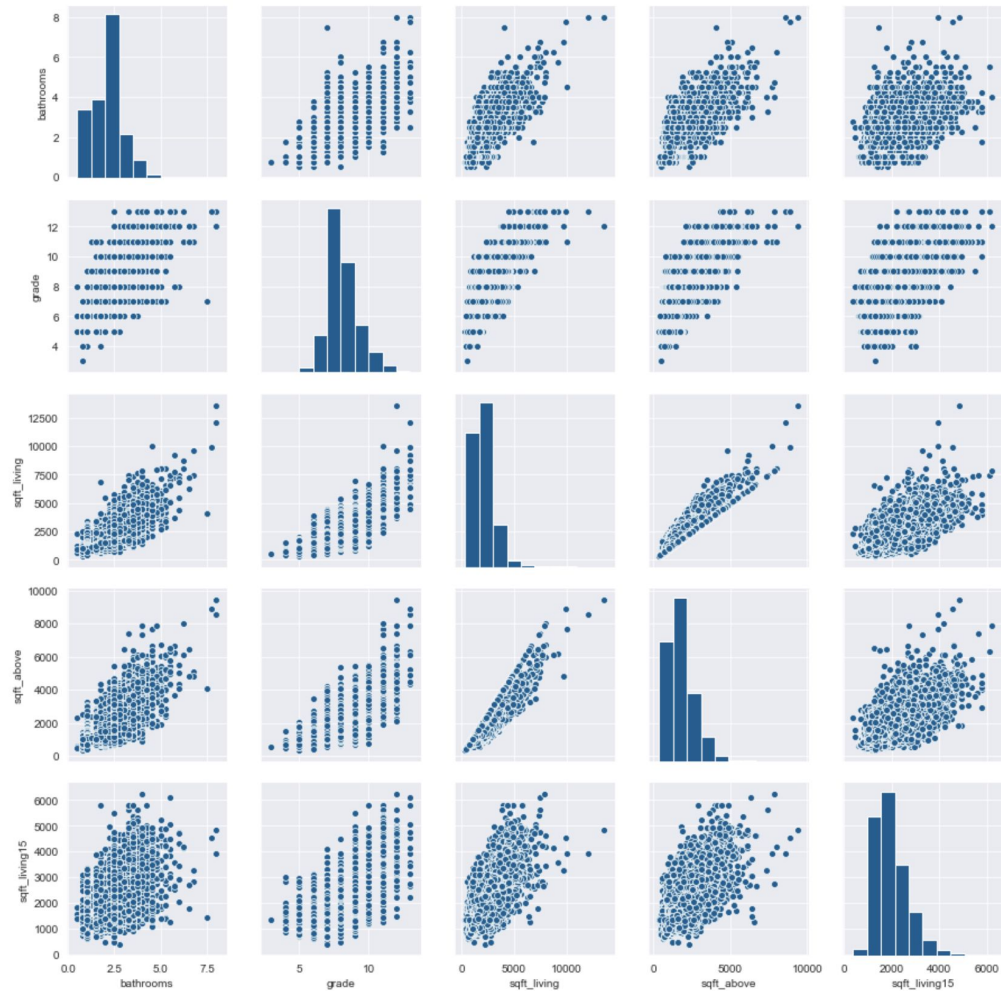
bedrooms_bin = disc_5k.fit_transform(df[['bedrooms']])
bathrooms_bin = disc_3.fit_transform(df[['bathrooms']])
```

Scaling factors:

```
from sklearn import preprocessing
from sklearn.preprocessing import RobustScaler, StandardScaler

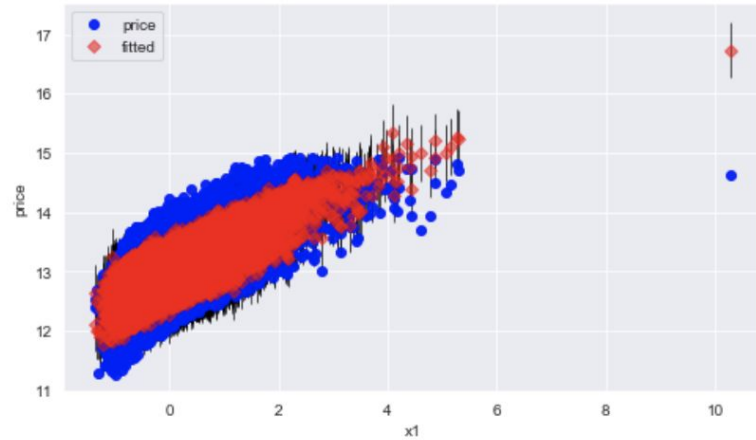
scale_robust = RobustScaler(copy=False, quantile_range=(25.0, 75.0), with_centering=True,
                             with_scaling=True)

scale_std = StandardScaler(copy=False)
```

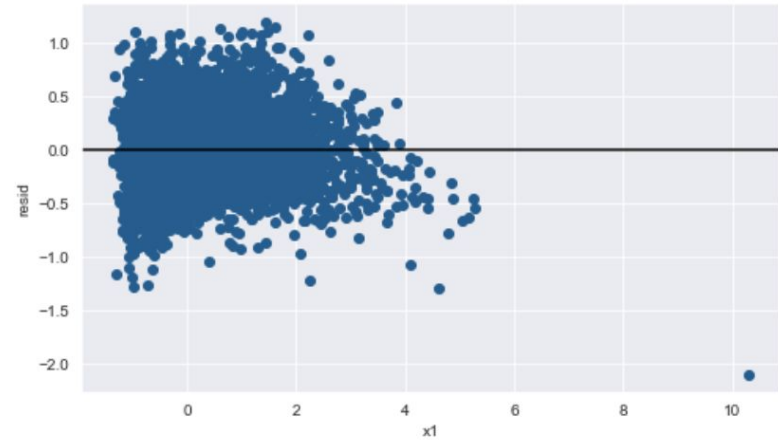


# Regression Plots for x1

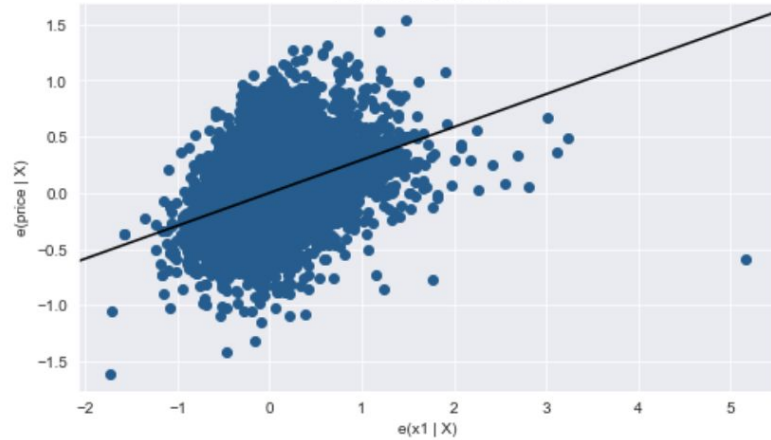
Y and Fitted vs. X



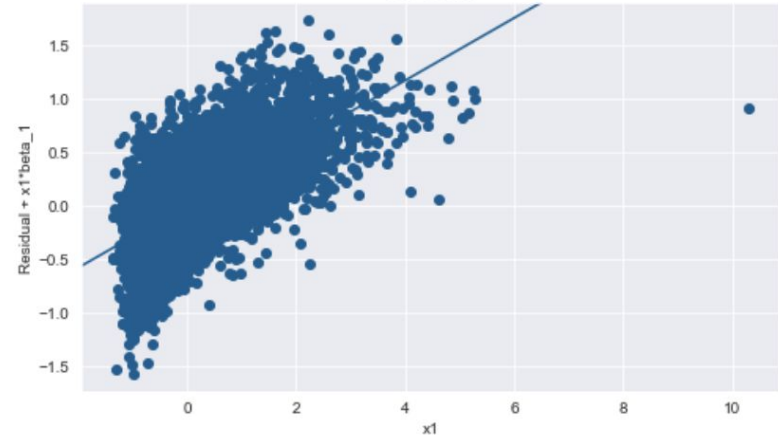
Residuals versus x1




Partial regression plot



CCPR Plot







<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.787
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.787
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2231.
<b>Date:</b>	Wed, 08 May 2019	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	10:50:44	<b>Log-Likelihood:</b>	130.69
<b>No. Observations:</b>	12068	<b>AIC:</b>	-219.4
<b>Df Residuals:</b>	12047	<b>BIC:</b>	-64.02
<b>Df Model:</b>	20		
<b>Covariance Type:</b>	nonrobust		

```
mymean = np.full((len(y_test), ), np.mean(y_train))
print("MSE: ", metrics.mean_squared_error(y_test, mymean))
print("RMSE: ", np.sqrt(metrics.mean_squared_error(y_test, mymean)))
rmse = np.sqrt(metrics.mean_squared_error(y_test, mymean))
```

```
MSE: 0.2660806452187228
RMSE: 0.5158300545903881
```

 <code>setup.py</code>	Require <code>scikit-learn&gt;=0.15.0</code> . <a href="#">Resolves #49</a> .	3 years ago
 <code>tox.ini</code>	added <code>pandas 0.22</code> and <code>sklearn 0.19</code> to <code>tox</code> .	a year ago

## README.rst

# Sklearn-pandas

 PASSED

This module provides a bridge between [Scikit-Learn](#)'s machine learning methods and [pandas](#)-style Data Frames.

In particular, it provides:

1. A way to map `DataFrame` columns to transformations, which are later recombined into features.
2. A compatibility shim for old `scikit-learn` versions to cross-validate a pipeline that takes a `pandas DataFrame` as input. This is only needed for `scikit-learn<0.16.0` (see [#11](#) for details). It is deprecated and will likely be dropped in `sklearn-pandas==2.0`.
3. A couple of special transformers that work well with pandas inputs: `CategoricalImputer` and `FunctionTransformer`.

## Installation

You can install `sklearn-pandas` with `pip`:

```
# pip install sklearn-pandas
```

## Tests

The examples in this file double as basic sanity tests. To run them, use `doctest`, which is included with python:

```
# python -m doctest README.rst
```





# Takeaways

- + My models were able to predict the price of a house with an adjusted R-squared of 0.79 and a MSE of 0.27.
- In this dataset, `sqft_living` and `grade` are the variables most highly correlated with price.
- R-squared and MSE/RMSE can diverge sharply, and multiple metrics of accuracy must be taken into account.
- The choice of which variables to keep and which to eliminate for multicollinearity has significant impacts on the accuracy of the final model.
- Data exploration and cleaning take up more time than running the actual regressions, but are crucial precursors to manipulating and processing the information.
- Pandas and scikitlearn contain many powerful tools with extensive documentation to aid in statistical analysis of data, but these libraries present a surprising number of incompatibilities that must be resolved via ad-hoc methods (or discovery of additional libraries like `sklearn-pandas`).